

DL for ENG

LECTURE 2.1 · RECORDED

Scientific Python


Remus Teodorescu

*With support from Research Assistant Noman Khan -
nomank@energy.aau.dk*

Aalborg University · ret@et.aau.dk

Deep Learning for Engineering · 2026

PANDAS · SOLVE_IVP · CURVE_FIT



An aluminium block cooling in a room, logged every two seconds. It arrives as a messy file, is solved as a differential equation, and is fitted as a two-parameter model. That fitted model is the baseline every network here has to beat.

Learning Objectives

SECTIONS, IN ORDER

- 1 **Pandas** — the DataFrame, reading a real file, selecting, missing data, aggregating
- 2 **Solving** — the state form, solve_ivp, and how to know the answer is right
- 3 **Fitting** — a model with two knobs, curve_fit, the covariance, and the baseline

Everything here sits on the arrays of the last lecture. Pandas adds labels and missing values, SciPy adds solvers, and underneath both it is still NumPy.

READ, AND WHAT EACH ANSWERS

pandas guide	how do I get a real file into an array?
SciPy integrate	how do I solve it, and is the answer right?
SciPy optimize	what are the parameters, and how well known?

Recap

1 – 2 · THE ARRAY

1 What is a NumPy array, and where does its speed come from?

One block of memory with one type and one shape. Arithmetic on it runs element by element in compiled code, with no Python loop, which is about a hundred times faster than a list. Writing the calculation that way is vectorisation. *NumPy basics, arrays; Ex_01 notebook 02*

2 What does the shape of an array tell you, and what habit goes with it?

The shape is a tuple with one length per axis: (3, 4) is three rows of four. Indexing with an integer removes an axis; a slice keeps it. The habit is to print the shape on the line after you make an array. *NumPy basics, shape, indexing*

3 – 4 · BROADCASTING

3 What is the broadcasting rule?

Write the two shapes right-aligned and compare them axis by axis: two lengths are compatible if they are equal or one of them is one, and an axis of one is stretched. A missing axis on the left counts as one. *NumPy basics, broadcasting*

4 Why is broadcasting dangerous, and how do you guard against it?

It raises nothing when two shapes are compatible by accident: a (100,) array minus a (100, 1) array gives (100, 100), ten thousand numbers, and a loss computed from it is wrong without a warning. Write the two shapes out, and check the shape of the result. *NumPy broadcasting; Ex_01 notebook 03*

DataFrame

DataFrame — a two-dimensional table with a name on every column and an index on every row. Each column is a NumPy array with its own dtype, so one table can hold a timestamp, a float channel and a status string. Think of it as a dictionary of columns rather than a list of rows, because that is how it is stored.

AND ONE COLUMN AT A TIME, NOT ONE ROW AT A TIME

The index is the part with no NumPy equivalent. It labels the rows, it need not be integers, and when it is a time it unlocks resampling and interpolation for nothing.

Everything from the last lecture still applies underneath. A column is an array, broadcasting works, and the shape rule is unchanged. Pandas adds labels and missing values on top.

THE OBJECT

index	time_s	temp_C	status
0	0.0	80.05	OK
1	2.0	78.11	OK
2	4.0	NaN	FAULT
3	6.0	74.90	OK

THREE THINGS TO CALL FIRST

```
df.head()      the first five rows
df.dtypes      one dtype per column
df.describe()  count, mean, min, max
# and df.shape, every time
```

Reading a CSV

read_csv — the one function that turns a delimited text file into a DataFrame. In the easy case that is the whole story. In the real case, where the file came off a logger, six of its arguments do almost all the work.

ONE FUNCTION DOES ALMOST ALL OF IT

sep, when the file is not comma-separated, because European exports are usually semicolons. skiprows, for the header block a logger writes. decimal, for the comma decimal separator.

na_values, for whatever the instrument writes when a channel drops out. parse_dates, for timestamp columns. And usecols, when the file has forty channels and you want three.

THE ARGUMENTS THAT MATTER

```
import pandas as pd

df = pd.read_csv(
    "cooling_log.csv",
    sep=";",          # not a comma
    decimal=".",      # European decimals
    skiprows=2,       # logger preamble
    na_values=["", "SENSOR FAULT"],
    usecols=["time", "temp_C"],
)

df.dtypes           # look, every time
```

An object dtype where you expected a float is a file problem, not a pandas problem.

Selecting

loc and iloc — the two indexers, and the difference matters. loc selects by label, and its slices include their endpoint, because a label has no arithmetic. iloc selects by position, and its slices exclude the endpoint exactly as an array's do.

TWO INDEXERS, AND WHICH IS WHICH

Mixing them up gives an off-by-one that does not raise. The rule worth adopting: iloc when you are thinking about positions, loc when you are thinking about names or times.

Boolean masks work as they did last lecture and are the right way to express a condition: the temperatures above fifty degrees, the rows where the status column says running.

THE FOUR FORMS

```
df["temp_C"]           a Series
df[["time_s", "temp_C"]] a DataFrame

df.iloc[0:5]           by position, [0,5)
df.loc[0:5]            by label,  [0,5]

df.loc[df["temp_C"] > 50.0]
df.loc[df["status"] == "OK", "temp_C"]

t = df["time_s"].to_numpy()
T = df["temp_C"].to_numpy()
```

The last two lines are the exit door. Everything after slide ten is NumPy.

Missing Data

Missing data — a value the instrument did not give you, held as not-a-number. It propagates, so any arithmetic involving it is also not-a-number. That is a feature: a missing value cannot quietly become a wrong one.

AND ONE DISHONEST ONE

There are three defensible responses, and which you choose depends on why the value is missing rather than on convenience. Drop the row, interpolate across a short gap, or carry the last reading forward.

The dishonest option is filling with zero. A missing temperature is not zero degrees, and a zero drags every mean, every fit and every plot towards itself while looking like data.

TALK · 2 MIN How would you decide what to do with a missing measurement, in your own words? Then, in the three options on the right, which one would you choose for a thermocouple that dropped out for two seconds, and what would you write in the report?

THE THREE

dropna()

the reading is absent — drop the row

interpolate()

smooth quantity, short gap

keep the NaN

use mean(), which skips it

```
df["temp_C"].isna().sum()      1
clean = df.dropna(subset=["temp_C"])
len(df), len(clean)           6, 5
```

THE ONE TO REFUSE

fillna(0.0). A missing temperature is not zero degrees.

Aggregating

groupby — an operation that splits a table by the values in a column, applies a reduction to each group, and puts the results back together. One line for what would otherwise be a loop, a dictionary and a concatenation.

TWO OPERATIONS THAT WOULD OTHERWISE BE LOOPS

In engineering the grouping column is usually a run identifier, an operating point or a sensor name. Six gauges, mean and standard deviation of each: one line.

resample is the same idea on a time index, and it is how a logger's irregular ten-hertz stream becomes a regular one-second series. That is all this course says about pandas.

ONE LINE EACH

```
# mean and spread per gauge
g = df.groupby("gauge")["strain"]
g.agg(["mean", "std", "count"])

# a regular 1 s series from an
# irregular log
reg = df.set_index("t").resample("1s").mean()

# resampling interpolates. That is a
# modelling choice, not formatting.
```

AND THAT IS ENOUGH PANDAS

You will need about five per cent of this library. Look the rest up on the day, and do not read the documentation cover to cover.

The State Form

The state form — the one form every numerical integrator accepts: a first-order system. You supply a state vector, an initial condition, and a function that takes the current time and state and returns the rate of change of each component.

A FIRST-ORDER SYSTEM, AND A FUNCTION FOR ITS RATES

A higher-order equation is reduced to it by making the derivatives part of the state. A second-order equation becomes a two-component system of position and velocity.

So the whole of the modelling work is writing the right-hand side function. It takes t and y , in that order, and returns the derivatives, even when it does not use t .

THE ONLY FORM A SOLVER ACCEPTS

$$\frac{dy}{dt} = f(t, y), \quad y(t_0) = y_0$$

$$\ddot{x} = g(t, x, \dot{x}) \implies y = (x, \dot{x}), \quad \dot{y} = (\dot{x}, g)$$

THE RIGHT-HAND SIDE, IN FULL

```
def cooling_rhs(t, y, tau, T_inf):
    """dT/dt in K/s. y = (T,)"""
    T, = y
    return [-(T - T_inf) / tau]

# t comes first, even when unused
```

solve_ivp

solve_ivp — SciPy's initial-value solver. It takes the right-hand side, the time span as a pair, the initial state, and the argument people forget, `t_eval`, which is the times you actually want back rather than the times the stepper happened to take.

FOUR ARGUMENTS THAT MATTER, AND TWO THAT DO NOT

The method defaults to an adaptive Runge–Kutta scheme, which is the right first choice for almost everything here. `rtol` and `atol` set the error it will accept, and the defaults are loose.

The result holds `t` and `y`, and `y` comes back as states by times. A single-state system therefore gives you shape one by `N`, which is the shape that will surprise you.

THE CALL

```
from scipy.integrate import solve_ivp

sol = solve_ivp(
    cooling_rhs,
    t_span=(0.0, 60.0),
    y0=[80.05],
    t_eval=np.linspace(0.0, 60.0, 301),
    args=(60.0, 22.0),
    method="RK45", rtol=1e-8, atol=1e-10,
)

sol.y.shape          (1, 301)
T_num = sol.y[0]     (301,)
```

The (1, 301) is a length-one axis. You met that trap last lecture.

Checking the Answer


Checking a solver — three checks, in decreasing order of availability. Compare against an analytic solution where one exists. Check a conserved quantity, which most engineering systems have. And halve the tolerance to see whether the answer moves.

THREE CHECKS, IN DECREASING ORDER OF AVAILABILITY

The third is the weakest and always available. If tightening the tolerance by two orders of magnitude changes your third significant figure, the answer you had was not converged.

This is not pedantry. From the eighth lecture onwards you validate physics-informed networks with these three, in this order, because there is often no test data at all.

NUMERICAL AGAINST ANALYTIC



Line: exponential. Dots: RK45. Max difference $\approx 1e-8$ °C.

THE THREE, RANKED

best	against an analytic case
usually	a conserved quantity
always	halve the tolerance, look again

The Model

Curve fitting — choosing the parameters of a model so that it best matches measurements, which is the same three ingredients as every model in this course: a function with parameters, a measure of how wrong it is, and something that makes that measure smaller.

THE SAME THREE INGREDIENTS AS EVERY MODEL HERE

The difference from a network is the function. Here it is a formula you wrote, with two or three parameters that mean something physical, rather than a few thousand that do not.


That is why a fit is the baseline a network has to beat, and why it is often the answer rather than the baseline.

$$\hat{\theta} = \operatorname{argmin}_{\theta} \sum_{k=1}^N (f(t_k, \theta) - T_k)^2$$

THE MODEL FUNCTION

TWO FREE PARAMETERS

```
def cooling_model(t, T0, tau):
    """T in °C. t in s. T_inf fixed."""
    return 22.0 + (T0 - 22.0) * np.exp(-t / tau)
```



Two numbers, chosen to make the vertical distances as small as possible.

curve_fit

curve_fit — SciPy's least-squares fitter. Give it your model function, the measured inputs and outputs, and a starting guess, and it returns the best parameters and their covariance matrix. The starting guess matters more than people expect.

THE ARGUMENT PEOPLE SKIP, THE OUTPUT THEY IGNORE

Pass `p0` whenever you can. Without it the optimiser starts from ones, which for a decay constant of ten to the minus four is far enough away to find nothing at all.

Pass `sigma` too when the measurements have different uncertainties, because otherwise every point is weighted the same and the noisiest one pulls hardest.

THE CALL, AND WHAT TO DO WITH IT

```
from scipy.optimize import curve_fit

popt, pcov = curve_fit(
    cooling_model, t, T,
    p0=[80.0, 63.7],      # from physics
)

T0_hat, tau_hat = popt
perr = np.sqrt(np.diag(pcov))

tau = 59.4 +/- 1.8 s
```

$$\sigma_i = \sqrt{C_{ii}}, \quad \theta_i \pm 1.96 \sigma_i$$

A parameter without an uncertainty is half a result.

The Covariance

The covariance — the second thing `curve_fit` returns, and the one people throw away. Its diagonal holds the variances of the fitted parameters, so the square root of each is that parameter's standard error, and the off-diagonal terms say which parameters trade off against each other.

FOUR THINGS, AND THEY ARE NOT SMALL

A parameter with a standard error as large as itself has not been determined by your data, however good the curve looks on the plot.

Report the value and its error together. A time constant of twelve seconds means nothing until it is twelve plus or minus one, or twelve plus or minus nine.

TWO AGAINST TWO THOUSAND

meaning	τ has units; w_{47} does not
uncertainty	sqrt of the covariance diagonal
extrapolation	the physics continues; a network does not
ratio	2 parameters, 31 samples

AND WHAT IT DOES NOT SAY

Flexibility. If the functional form is wrong, no amount of fitting rescues it — and that is precisely the case a network is for.

The Baseline

The baseline — the fitted physical model that any network you build later has to beat. It has few parameters, each with a meaning and an error bar, it extrapolates as well as its physics does, and it takes a second to fit.

THREE MODELS, ONE DATASET, A RANKING THAT REVERSES

When it fails, it fails visibly: the residual shows structure, the covariance blows up, or the optimiser does not converge. All three are easier to notice than a network being quietly wrong.

So fit the physical model first, always, and keep its number. From the third exercise onwards that number is what every later result is compared against.

TALK · 2 MIN What does a fitted physical model give you that a network does not? Then, in the plot on the right, what would a network have to do before you would believe it over that curve?

THE MEASURED NUMBERS


A · analytic 1.12 °C → 0.74 °C

B · curve fit 0.32 °C → 0.05 °C

C · network 0.12 °C → tens of °C

training RMSE → extrapolation RMSE

WHY THE NETWORK LEAVES



Advanced Topics

RELATED TOPICS LEFT OUT, AND WHERE THEY MATTER

topic	what it adds	read	where it makes a difference
merge and join	two tables keyed on a column	pandas guide, merging	measurements against a run log
stiff solvers	implicit steppers for fast modes	SciPy Integrate, ODE	chemistry, circuits, control
bounded optimisation	parameters kept physical	SciPy optimize, least_squares	a rate that cannot be negative
interpolation	values between the samples	SciPy interpolate	resampling, lookup tables
signal processing	filters and spectra	SciPy signal	vibration, power quality

WHY THEY ARE LEFT OUT

None is needed for the exercises, and each is a chapter of its own. Two will find you anyway: a stiff solver the first time a solution oscillates for no reason, and bounds the first time a fit returns a negative time constant.

Exercise 2

AND WHICH NOTEBOOK ACTUALLY COVERS THIS

There is no separate exercise notebook for this half-lecture, and I would rather say so than pretend otherwise. Ex02 belongs to lecture two point two and is entirely PyTorch.

The material here is exercised in Ex03, in three weeks, where you fit the cooling curve by hand and then compare it to a network. That is the right place for it, because the comparison is the point and it needs a network to compare against.

So there are two things to do now. Finish Ex01 notebooks 02 and 03 if you have not, because lecture two point two assumes both. And spend twenty minutes typing the code blocks from today's slides into a notebook of your own against any CSV you have — a lab log, a spreadsheet export, anything.

That second one is a genuine recommendation rather than a filler task. Reading a file you did not choose is the only way to meet the six problems on slide seven, and every one of them will be in there somewhere.

WHERE THIS MATERIAL IS EXERCISED

Ex01 02–03 finish these — L2.2 assumes them

Ex02 next lecture, PyTorch, all of it

Ex03 01 the curve fit, by hand

Ex03 03 and the comparison that reverses.
Ex03 fits the same block again, so the number carries over.

REPORT QUESTION, TO CONCLUDE

Across the notebooks: give the fitted time constant with its error bar, say which of the three solver checks you ran, and state how many rows you dropped and why.

Key Takeaways

MEASUREMENTS

what it is	a table of labelled columns, arrays underneath
function	reads a real file, selects, groups, handles gaps
train, tune	check dtypes after every read; say what you dropped
vs a bare array	labels, a time index, and missing values
used for	every measurement that arrives as a file

SOLVING AND FITTING

what it is	an integrator and a least-squares fitter
function	state in, trajectory out; data in, parameters out
train, tune	t_eval and the tolerance; p0 and sigma
vs a network	few parameters, each with a meaning and an error
used for	the baseline every later model has to beat

Questions

1 – 2 · MEASUREMENTS

1 What is a DataFrame, and what does it give you that an array does not?

A table of labelled columns with arrays underneath. Each column has its own type, the rows carry an index, often time, and a missing value is marked as missing. So you select by name and by time, not by position.

pandas guide, intro; Ex_02 notebook 01

2 What do you check after reading a measurement file, and how is a missing value handled?

The type of every column, the index and the number of rows, after every read. A missing value is dropped, interpolated or left marked, and the report says which and how many. Filling it with a number that looks measured is the dishonest answer. *pandas guide, IO, missing data; Ex_02 notebook 01*

3 – 4 · SOLVING AND FITTING

3 What form does an integrator demand, and how is a numerical solution checked?

A first-order system: the state in, its time derivative out, so a second-order equation becomes two first-order ones. Check against an analytic solution where one exists, against a conserved quantity, and, always available, by tightening the tolerance to see whether the answer moves. *SciPy integrate; Ex_02 notebook 02*

4 What does a least-squares fit return, and why is it the baseline for every later model?

The parameters that bring a model closest to the data, and a covariance matrix that gives each one an error. The model has few parameters, each with a physical meaning. A parameter whose error is as large as itself is not determined by the data. *SciPy optimize; Ex_02 notebook 03*

NEXT

L2.2 — PyTorch as an Array Library

Today you differentiated nothing and solved everything by conventional means. Next time: arrays that remember what happened to them, and a differentiation engine that will still be doing the work in week twelve. It is the most load-bearing half-lecture in Part 1.

BEFORE THE NEXT RECORDING

finish	Ex01 notebooks 02 and 03
try	pandas on a CSV of your own
read	the solve_ivp worked example
expect	Ex02, and to spend real time on it

*Lecture two point two is the one to watch twice.
Everything in Part 2 depends on one function call in it,
and the exercise notebook that goes with it is the
most important in Part 1.*