

DL for ENG

LECTURE 3.1

What is Artificial Intelligence?


Remus Teodorescu

*With support from Research Assistant Noman Khan -
nomank@energy.aau.dk*

Aalborg University · ret@et.aau.dk

Deep Learning for Engineering · 2026

DEFINITIONS · HOW IT LEARNS · JUDGEMENT



Three words that are used as if they meant the same thing, and do not: artificial intelligence, machine learning, deep learning. Today is what each one is, where the training signal comes from in the three kinds of learning, and the judgement an engineer needs most, which is knowing when a trained network is the wrong tool and a solver is the right one.

Learning Objectives

SECTIONS, IN ORDER

- 1 Definitions** — artificial intelligence, machine learning, the neural network, deep learning, the language model
- 2 How it learns** — supervised, unsupervised, reinforcement
- 3 Judgement** — the modelling spectrum, where it fits, when not to use it, how much data

No mathematics today beyond one energy function and one line defining optimisation. The machinery starts in the next lecture.


READ, AND WHAT EACH ANSWERS

- | | |
|--------------------------|---|
| UDL ch. 1 | what do the four words actually mean? |
| UDL ch. 1.1 – 1.3 | where does the training signal come from? |
| UDL ch. 1.4 | when is a network the right tool, and when not? |

Artificial Intelligence

Artificial intelligence — the whole ambition of making a machine do what would take intelligence in a person: search, planning, logic, robotics and language. Most of it involves no learning at all. Machine learning is the part that fits models to data, and deep learning is the part of that which stacks many layers.

WHO EXTRACTS THE FEATURES



FOUR AREAS, ONE INSIDE THE NEXT

ARTIFICIAL INTELLIGENCE

search · planning · reasoning · robotics

MACHINE LEARNING

regression · k-means · PCA · SVM · trees

NEURAL NETWORKS

perceptron · MLP · Hopfield · Boltzmann

DEEP LEARNING

CNN · RNN · GNN · transformers · PINNs


THE ONE TO REMEMBER

A Kalman filter is AI by any 1970s definition, and still the best tool for state estimation in most control rooms, where it is still the baseline a neural method has to beat, and often will not.

Machine Learning

Machine learning — fitting a model to data so that it decides or predicts without being given the rule. It takes three ingredients: a function with adjustable parameters, a loss that says in one number how wrong it is, and an optimiser that makes that number smaller. Linear regression counts, and so does a billion-parameter network.

THE THREE INGREDIENTS



$$\hat{\theta} = \operatorname{argmin}_{\theta} \mathcal{L}(\theta)$$

THE SAME THREE, LECTURE BY LECTURE

L4	the function gets deep
L5	the function gets structure
L6.1	the loss gets built properly
L7	the loss gets physics in it
L12	the function is a whole power grid

WHAT ACTUALLY CHANGES IN PART 2

Only the middle one. The function stays a network and the optimiser stays gradient descent. Physics-informed learning is a statement about the loss, and nothing else.

THE THREE INGREDIENTS IN PYTORCH

```

model = nn.Linear(1, 1)           # the function: two knobs
loss = ((model(x) - y) ** 2).mean() # the loss: how wrong it is
opt = torch.optim.SGD(model.parameters(), lr=0.1) # the optimiser
opt.zero_grad(); loss.backward(); opt.step()      # better knobs
  
```

Neural Network

Neural network — a function built from many small units, each taking a weighted sum of its inputs and passing it through a nonlinearity. The weights are what training adjusts. It was written down in 1943 as a model of a biological neuron, and the borrowed word has misled people ever since.

FROM THE FIRST NEURON TO THE NOBEL PRIZE



- **1943 • McCulloch and Pitts** — a threshold on a weighted sum
- **1958 • Rosenblatt** — the perceptron, built in hardware
- **1969 • Minsky and Papert** — one layer cannot do XOR
- 1986 • backpropagation trains many layers, ending the winter
- **2024 • Hopfield and Hinton** — the Nobel Prize in Physics

HOPFIELD'S ENERGY

$$E = -\frac{1}{2} \sum_{i,j} w_{ij} s_i s_j - \sum_i \theta_i s_i$$

Neural networks began in statistical physics: the Hopfield network (1982) and Hinton's Boltzmann machine (1985) are spin-glass energy landscapes. This course takes them back to engineering; the prize that recognised the origin is two slides on.

A REAL NEURON, AND ITS COPY



Deep Learning

Deep learning — machine learning with a function made of many layers, each one feeding the next, fitted end to end by gradient descent. What the layers buy is that nobody designs the features: the early layers learn what to measure and the later ones learn what to do with it. The method was correct in 1986 and unusable until the data and the arithmetic arrived.

WHY 2012 AND NOT 1992

- **1986 to 2012** — the mathematics barely changed, and SVMs won
- **2012 · AlexNet** — 1.2 million labels, two graphics cards
- **the lesson** — too little data makes a correct method look wrong

ALEXNET, 2012 · KRIZHEVSKY, SUTSKEVER AND HINTON



Eight layers, 60M parameters. 15.3 % ImageNet error, against 26.2 %.


WHAT ARRIVED, AND WHEN

the idea	1986 — and then it waited
the data	ImageNet, 2009, 14 million images
the arithmetic	GPUs, and someone willing to try


THE ONE TO REMEMBER

Too little data or compute makes a correct method look wrong.

GPU VS CPU



CPU · 8 cores, in sequence



GPU · thousands, all at once

A matrix product is one operation on thousands of numbers. CUDA (2006) let 3D-game cores do any arithmetic.

Large Language Model

Large language model — a network trained on an enormous quantity of text to predict the next word, then trained further to follow instructions. What it computes at each step is a distribution over the next piece of text, and nothing else. It is strong where the answer has been written down, and unreliable where it must be derived.

WHAT THE BENCHMARKS ACTUALLY ASK


- **exam level** — a PhD scores 65 % on GPQA, models passed it
- **two meanings of research** — find one, or derive one
- in engineering the check is physical

WHAT 'RESEARCH' MEANS ON A LEADERBOARD

#	model	maker	score	what it blends
1	GPT-6 Astra	OpenAI	93.8	GPQA · HLE ·
2	GPT-5.6 Sol	OpenAI	93.4	Browse Science
3	GPT-5.6 Terra	OpenAI	90.2	DeepSearch-QA
9	Claude Mythos 5	Anthropic	82.2	GAIA
				benchlm.ai, 76 models

Every benchmark in the blend has an answer that exists; CritPt's has to be derived. Knowing what is known is 93.8; working out what is not is 5.7.

EXAMS PASSED. RESEARCH NEXT.



THE CONCLUSION


Language models accelerate your research; they do not replace you in it. A language model is fluent about physics; a physics-informed network is constrained by it. The PINN on our oscillator is wrong by 0.59 mm and knows which equation it is solving. Use the first to build the second.

Supervised Learning

Supervised learning — learning from inputs whose answers are already known: every example carries the output it should produce, and the answers are the supervision. It is the setting of almost everything in Part 1. Its cost is that somebody had to produce those answers, by measuring or by labelling, before any training could start.

WHERE THE ANSWERS COME FROM

- **the answers are the supervision** — somebody made them
- **in engineering** — a simulation is a label factory, run once
- **Ex_03 is supervised** — time in, displacement out



the answers already exist

IN THIS COURSE

Ex_03	fit a curve, three ways
Ex_04	classify, and fail at XOR
Ex_05	images, and a six-bus grid
L11.1	steering angle from a camera

TALK · 2 MIN Where would the labelled answers for your own project come from, and who would have to produce them? Then, in the table on the right, which of those four is the one where labels are cheapest, and why?


Unsupervised Learning

Unsupervised learning — learning from data with no answers attached, to find the structure that is already in it: clusters, axes of variation, or departures from normal. Its other face is learning features from unlabelled data, which is how a model is pre-trained before a handful of labels are used. Nothing guarantees the structure it finds is the one you wanted.

STRUCTURE, AND FEATURES WITHOUT LABELS

- **the structure already there** — clusters, axes, anomalies
- **in engineering** — learn normal from logs, then flag departures
- **the catch** — the structure found need not be the one you wanted

GOOGLE BRAIN 2012 · ONE NEURON FIRES FOR CATS



WHAT IS IN SCOPE

PCA	yes — mentioned in L6.2
clustering	yes — briefly
anomaly detection	yes — the engineering case
generative models	out of scope — UDL ch. 14 to 18

LEARNING FEATURES WITHOUT LABELS

Hinton, 2006	deep belief nets: unlabelled digits, a few labels on top
Google Brain, 2012	a 'cat' neuron from 10 million unlabelled frames
language models	pre-trained the same way, on unlabelled text


Reinforcement Learning

Reinforcement learning — learning to act from consequences rather than answers. The model acts, the world responds with a reward that usually arrives late, and the model adjusts to make high-reward actions more likely. Supervised learning is told the answer, unsupervised learning is told nothing, and this is told how good its answer was, afterwards.


ACT, OBSERVE, ADJUST

- **no answers, only consequences** — and they arrive late
- the reward is the design problem, not a detail
- **the cost** — millions of trials, so it lives in a simulator

LEARNING TO WALK · TRY, FALL, TRY AGAIN



attempt → consequence → adjust. No labels, no dataset: the floor is the teacher.



reward against attempts

THE VOCABULARY, TRANSLATED

state	what the plant is doing
action	the control input
policy	the controller
reward	minus the cost function
episode	one run of the experiment

WHERE IT IS TAUGHT

L6.2 teaches the fundamentals and the role of reinforcement in post-training. L11 runs it on a JetRacer, where a bad policy costs you a physical crash rather than a number in a table.

The Modelling Spectrum

The modelling spectrum — the line running from a model built entirely from physics to one built entirely from data, on which every model sits somewhere. Pure physics extrapolates and is slow. Pure data is fast and silent outside the range it was fitted on. The question is never physics or data, but how much of each.

ONE LINE, NOT TWO CAMPS

- **pure physics** — extrapolates, and is slow
- **pure data** — fast, and silent outside its range
- every model you have built already sits somewhere on it



THE LINE, END TO END



pure physics	no data, extrapolates, slow
hybrid	physics written into the loss — Part 2
pure data	fast, silent outside its range

TALK · 2 MIN Take a model you have built yourself. Where on this line does it sit, and what would move it one step towards physics? Then, on the bar above, which of the five markers is closest to it?

Where It Fits

Where it fits — the four situations in which a trained network earns its place: interpolating where the data is dense, working in many dimensions where nobody can write the correlations down, evaluating in microseconds once trained, and recovering a parameter nobody can measure.

FOUR SITUATIONS, ONE SHAPE

- **interpolation** — where the data is dense, Ex_03 gives 1.8 mm
- **many dimensions** — correlations nobody can write down
- **speed** — microseconds an evaluation, once trained
- **inverse problems** — a parameter you cannot measure directly

THE COMMON SHAPE

Complicated function, simple objective — that is the shape.

The method supplies the function. You must supply the objective.

Which means the engineering judgement moves entirely into the loss. Choose it badly and you will get a model that optimises exactly what you asked for and nothing you wanted.

In Part 2 the loss stops being a formality and becomes the design.

When Not To Use It

When not to use it — the three conditions that should stop you: data too thin to support the parameters, a guarantee that an error measured on the data you happened to test cannot give, and physics that is already known and cheap to solve.

YOUR DATA IS THIN

Fit a curve instead. A model with a handful of physical parameters is defensible on the data you have; one with thousands is not, however good the training error looks.

YOU NEED A GUARANTEE

Certification wants bounds. A trained network gives an empirical error on the data you happened to test — a different kind of object.

THE PHYSICS IS KNOWN AND CHEAP

If a solver runs in a second, use the solver. A network that took six hours and is wrong unpredictably is not an improvement.

ASK IN THIS ORDER

- | | |
|--------------|---|
| first | can I write the equation down? |
| then | is solving it fast enough? |
| then | do I have data in the hundreds or the millions? |
| then | who has to believe the answer? |
| last | and only then, which architecture? |

Almost every failed machine-learning project in engineering skipped straight to the last question. The ordering is the content of this slide.

Data Requirements

Data requirements — how many examples a model needs, which has no answer on its own. It depends on how much physics you write into the model: with none, thousands and still nothing trustworthy outside them; with all of it, no measurements at all. How much data is the wrong question asked first.

THE SAME PROBLEM, THREE WAYS

- **no physics** — thousands of examples, and still blind outside them
- **some physics** — 101 noisy points, 0.59 mm outside
- **all physics** — the equation and one initial condition

ORDERS OF MAGNITUDE

a thesis test campaign	10^2 runs
a year of plant SCADA	10^6 samples, few events
ImageNet	10^6 labelled images
a modern language model	10^{12} tokens
a PINN, in Part 2	zero labelled outputs

The last row is not a trick. In Part 2 you will train a network that has never been shown a correct answer, and it will produce one.

Advanced Topics

RELATED TOPICS LEFT OUT, AND WHERE THEY MATTER

topic	what it adds	read	where it makes a difference
symbolic AI	logic and search, with no learning	UDL ch. 1 notes	planning, scheduling, proof
classical machine learning	SVM, trees and k-NN on small data	UDL ch. 1 notes	hundreds of rows, tabular data
generative models	a model that produces new data	UDL ch. 14 – 18	design search, synthetic data
Bayesian methods	a distribution over the answer	UDL ch. 1 notes	uncertainty you have to report
reproducibility	seeds, splits and honest reporting	UDL ch. 1 notes	every claim you will read

WHY THEY ARE LEFT OUT

Each is a field beside this one rather than a step inside it, and most pay off where the data is small or the answer has to carry a probability. This course follows the one method that runs through all twelve lectures.

Exercise 3

WHAT THE NOTEBOOK DOES

One notebook, one problem: a mass on a spring, measured for five seconds by a sensor that is noisy and was never zeroed.

You build the same system five ways, from the analytic solution through a fitted regression and a plain network to a network with the equation of motion in its loss, and you compare them on data they never saw.

The network agrees with the measurements to 0.62 mm. The measurements themselves are wrong by 1.80 mm, so it fitted the noise, and outside the window it is fifty times worse than the model that was told the physics.

The fifth model uses no measurements at all. The equation and one initial condition are enough.

THE QUESTION IT ANSWERS

Ex_03 which model would you sign?

models 1 – 3 what does fitting noise cost?

models 4 – 5 what does the equation buy?

REPORT QUESTION, TO CONCLUDE

All five models fit the data. Which one would you put in a design report, and what number in your own results defends that choice? The marks are for the reasoning, not for the error.

Key Takeaways

DEEP LEARNING

what it is	many layers, fitted end to end by gradient descent
function	inputs in, a prediction out, parameters fitted to data
train, tune	a function, a loss, an optimiser: the three ingredients
vs physics	faster, and blind outside the data it saw
used for	interpolation, many dimensions, speed, inverse problems

THE MODELLING SPECTRUM

what it is	the line from pure physics to pure data
function	places any model, including the ones you built
train, tune	the physics you write in decides the data you need
vs pure data	Ex_03: 0.59 mm outside the window against 30.26
used for	deciding whether to use a network at all

Questions

1 – 5 · DEEP LEARNING

- 1 What is artificial intelligence, and how are machine learning and deep learning nested inside it? *UDL ch. 1*
- 2 What does a machine-learning model compute, and what are the three ingredients it is built from? *UDL ch. 2*
- 3 What is a neural network, why did it take from 1986 to 2012 to work, and what arrived in between? *UDL ch. 1*
- 4 How does deep learning differ from classical machine learning in who extracts the features? *UDL ch. 1*
- 5 What does a large language model compute at each step, and what does that make it unreliable at? *UDL ch. 12*

6 – 10 · THE MODELLING SPECTRUM

- 6 What is the modelling spectrum, and where does a fitted correlation sit on it? *UDL ch. 1.4*
- 7 Where does the training signal come from in supervised, unsupervised and reinforcement learning? *UDL ch. 1.1 – 1.3*
- 8 In which four situations does a trained network earn its place, and what shape do they share? *UDL ch. 1.4*
- 9 What are the three conditions that should stop you using one, and which is hardest to argue with? *UDL ch. 1.4*
- 10 How much data does a model need, and what did Ex_03 measure outside the window for each model? *Ex_03; UDL ch. 1.4*

NEXT

L3.2 — What the Field Has Built

Today was what the field is and how it got here. Next time is what it has actually produced — across vision, language, control, and then across science and engineering, which is where this course is going.

BEFORE THE NEXT SESSION

read	UDL chapter 1 — about twenty minutes
run	Ex_03, to check your setup
bring	one model from your own work
think about	where on the spectrum it sits

Bring the model. L3.2 ends by placing the room's own examples on the spectrum, and a lecture theatre full of engineers produces better cases than any set I could invent — but only if you arrive having thought about one.