


DL for ENG

LECTURE 7.1

Physics-Informed Neural Networks



Mathematical foundations — formulation, loss construction and training

Remus Teodorescu

With support from Research Assistant Noman Khan - nomank@energy.aau.dk

Aalborg University · ret@et.aau.dk · Deep Learning for Engineering · 2026

Liu (2025), PINN with Python · Raissi, Perdikaris & Karniadakis (2019) · Prince, Understanding Deep Learning

What You Will Be Able To Do

THE GOALS, IN PLAIN WORDS

- **write** — a PDE residual with autograd, and check it against an analytic derivative
- **assemble** — a composite loss, and say what each term is for
- **explain** — why the activation must be smooth, from the order of the residual
- **enforce** — a boundary condition softly and exactly, and state what each costs
- **choose** — between collocation, finite differences and the energy form
- **say** — why 105 sample points can constrain 2,751 parameters

WHAT TODAY ASSUMES

- | | |
|------------|-----------------------------------|
| from L6.1 | a loss term need contain no data |
| from Ex_03 | a differential equation in a loss |
| from L4.1 | why tanh, and not ReLU |


YOU HAVE DONE THIS ONCE

Ex_03 put a damped oscillator in a loss and measured what it bought. The residual was given to you. Today you write it.

AND ON A PDE

Two independent variables instead of one, and boundary conditions instead of initial ones.


The Logical Flow of This Lecture



- **five stages** — each a mathematical object, not an implementation detail
- the network is an ansatz, autograd an exact operator, the loss a functional
- **L7.2 reuses this pipeline** — only the operator and the conditions change

Where L7 Sits: One Method, Five Applications

- **energy systems are governed by PDEs** — expensive to solve repeatedly
- and repetition is what design, control and optimisation require
- **every lecture that follows** — keeps this skeleton, changes only the physics




The framing for this course: physics enters through the loss, not the architecture — which is why the same network, unchanged, solves heat conduction in L8 and a battery thermal model in a project.

Why Embed Physics in the Learning Problem?

- **classical solvers discretise first** — accuracy inherited from the mesh
- **their unknowns are nodal values** — cost grows exponentially with dimension
- a PINN makes the solution a differentiable function of the coordinates
- **so the PDE can be tested anywhere** — no connectivity required
- **and the physics is imposed, not learned** — which is why little data suffices

network + automatic differentiation + PDE residual = mesh-free solver



KEEP THE CLAIM HONEST

PINNs are not a drop-in replacement for FEM. On stiff, high-frequency or multi-scale problems classical solvers remain far more reliable — see the failure modes later in this lecture.

Where PINNs Sit: The Modelling Spectrum

- ▶ PINNs are best understood not as a new kind of solver but as a bridge between two established traditions — and the loss function is the bridge.

Pure physics

FEM · FDM · spectral

Solve the governing equations by discretisation.
No data required — but a mesh is.
Accuracy inherited from the mesh.

PINN

physics in the loss

Equations and data constrain one continuous solution.
Mesh-free and differentiable.
Physics is a soft or hard constraint.

Pure data

ML surrogates

Fit observations with no physical constraint.
Fast to evaluate — but free to violate known laws.

no data, all structure

all data, no structure

This is why L7 spends its time on the loss function: the loss is where the physics actually lives.


The Network as a Solution Ansatz

- the unknown field becomes a network with trainable parameters
- inputs are the independent variables: x , t and any parameters
- each layer is an affine map then a smooth nonlinearity
- **differentiable to arbitrary order in its inputs** — which is the point
- nothing is problem-specific yet: the physics enters through the loss

ANSATZ

$$u(\mathbf{x}, t, \mathbf{c}) = \mathcal{N}_{\hat{\mathbf{w}}}(\mathbf{x}, t, \mathbf{c})$$

FEEDFORWARD ARCHITECTURE



Universal Approximation Theorem


STATEMENT

$$\sup_{\mathbf{x} \in K} |u(\mathbf{x}) - \sum_{j=1}^N a_j \sigma(\mathbf{w}_j^T \mathbf{x} + b_j)| < \varepsilon$$

- one hidden layer of finite width approximates any continuous function
- that is what licenses using a network as a solution surrogate at all
- **an existence result only** — silent on size, and on findability
- **two versions** — width (Cybenko, Hornik) and depth
- **depth is preferred in practice** — same accuracy, far fewer parameters



WIDTH → ACCURACY



wider network → uniformly smaller error

Network Capacity and the Sampling Condition

PSEUDO-DIMENSION OF THE AFFINE SPACE

$$P^* = (p + 1) + (N_n + 1) N_L$$

SAMPLES-NEURONS CONDITION

$$N_{\text{sample}} \geq P^* \quad N_{\text{sample}} \gg P^*$$

- **P^* counts neurons, not weights** — the pseudo-dimension of the affine space
- fewer than P^* independent sample points and the fit is under-determined
- **boundary and initial points count as extra** — the condition is on the interior set
- **note the asymmetry** — 105 sample points against 2,751 parameters
- it is the residual, not the data, that constrains the fit

WORKED EXAMPLE


NETWORK	P^*	PARAMS USED FOR
50×2	105	2 751 Poisson, unit square
100×2	206	10 601 transient heat, wave
200×4	808	121 601 plate with hole, NBCs

P^ and the parameter count are different quantities — they differ here by factors of 26 to 150. The sampling rule uses P^* .*

Why the Activation Must Be Smooth

- a residual of order k needs the network k -times differentiable
- **the activation is the only nonlinearity** — so it sets the ceiling
- **tanh is the default** — bounded, monotone, infinitely differentiable
- **ReLU is piecewise linear** — second derivative zero
- smoothness is a requirement here, not a stylistic preference

$\Sigma, \Sigma' \text{ AND } \Sigma'' \text{ FOR TANH}$



ACTIVATION	$\sigma(z)$	$\sigma''(z)$	SECOND-ORDER PDE?
tanh	$\tanh(z)$	$-2 \tanh z (1 - \tanh^2 z)$	Yes — the default
sin / SIREN	$\sin(\omega_0 z)$	$-\omega_0^2 \sin(\omega_0 z)$	Yes — best for oscillatory
sigmoid	$1 / (1 + e^{-z})$	$\sigma(1-\sigma)(1-2\sigma)$	Yes — but saturates
ReLU	$\max(0, z)$	$0 \text{ almost everywhere}$	No — residual collapses


Automatic Differentiation (Autograd)

- autograd applies the chain rule through the forward graph
- the derivatives are exact to machine precision
- no step size, no truncation error, no discretisation of the domain
- the mechanism that gives the optimiser gradients gives du/dx
- `create_graph=True` keeps it differentiable: second derivatives

RESIDUAL DERIVATIVES, IN THE BOOK'S STYLE

```
u = pinn(xyt)
grads = torch.autograd.grad(
    outputs=u, inputs=xyt,
    grad_outputs=torch.ones_like(u),
    create_graph=True)[0]
```

COMPUTATIONAL GRAPH



Autograd vs. Numerical and Symbolic Differentiation

SYMBOLIC

Manipulates expressions

analytically.

Exact, but expression swell:
derivative expressions grow
combinatorially with order.
Impractical for deep graphs.

NUMERICAL (FD)

Finite differences on a step h .


Truncation error $O(h^2)$ plus
round-off error $O(\epsilon/h)$ — the
two conflict, so accuracy is
bounded by an optimal h .

AUTOGRAD

Chain rule on the graph.

Exact to machine precision.
No step size to tune.
Cost scales with graph depth.
The PINN default.

FINITE-DIFFERENCE ERROR VS. STEP SIZE



- finite differences cannot be made arbitrarily accurate
- shrink h and truncation falls while round-off grows
- autograd has no such floor
- so residuals are evaluated on the graph, not a stencil

The General PDE Operator Form

SYSTEM EQUATION

$$\mathcal{F}(u; \mathbf{x}, t, \mathbf{c}) = 0, \quad \mathbf{x} \in \Omega, \quad t \in [0, t_{\text{end}}]$$

- **F is a differential operator** — linear or nonlinear, acting on u
- u depends on x, t and the physical parameters c
- Ω is the spatial domain; for a static problem time drops out

ELLIPTIC

Poisson · L8

$$\rightarrow \mathcal{F} = u_{xx} + u_{yy} + f$$

PARABOLIC

transient heat · L8

$$\rightarrow \mathcal{F} = u_t - \alpha(u_{xx} + u_{yy})$$


HYPERBOLIC

wave · L7.2

$$\rightarrow \mathcal{F} = u_{tt} - c^2(u_{xx} + u_{yy})$$

one operator, three equations — only $\mathcal{F}$ and the admissible conditions change

DOMAIN AND ITS BOUNDARY



Admissible Conditions: DBC, NBC, IC and IV

DIRICHLET (DBC)

$$u(\mathbf{x}, t) = \bar{u}(\mathbf{x}, t), \quad \forall \mathbf{x} \in \partial\Omega_{DBC}$$

u prescribed on part of the boundary.

NEUMANN (NBC)

$$\mathcal{B}(\mathbf{x}, t, \mathbf{c}) = 0, \quad \forall \mathbf{x} \in \partial\Omega_{NBC}$$

A derivative — flux, traction, strain — prescribed on the boundary.

INITIAL CONDITION (IC)

$$\mathcal{I}(\mathbf{x}, 0, \mathbf{c}) = 0, \quad \forall \mathbf{x} \in \Omega$$

The state of the field at $t = 0$.

INITIAL VELOCITY (IV)

$$\mathcal{V}(\mathbf{x}, 0, \mathbf{c}) = 0, \quad \forall \mathbf{x} \in \Omega$$

The time derivative at $t = 0$; required for second-order-in-time PDEs.

$\mathcal{F}$, $\mathcal{B}$, $\mathcal{I}$ and $\mathcal{V}$ are all evaluated through the same network by autograd — the operators differ, the machinery does not.

Forward and Inverse Problems

FORWARD PROBLEM

- ✓ operator $\mathcal{F}$ and parameters c
- ✓ boundary / initial conditions

INVERSE PROBLEM

- ✓ sparse measurements of u
- ✓ form of the operator $\mathcal{F}$

FORWARD — the usual direction



INVERSE — a parameter becomes unknown



WHY THIS IS WORTH HAVING — A BATTERY THERMAL EXAMPLE

A cell is instrumented with three surface thermocouples. The internal heat-generation rate and the thermal conductivity cannot be measured directly — but they are exactly the parameters c in the heat equation. Making them trainable recovers both from those three sensors, using the same network and the same loss. The classical alternative is to guess a value, run a forward solve, compare, and repeat.

The Composite Loss Functional

$$\mathcal{L} = \mathcal{L}_{data} + w_{PDE}\mathcal{L}_{PDE} + w_{DBC}\mathcal{L}_{DBC} + w_{NBC}\mathcal{L}_{NBC} + w_{IC}\mathcal{L}_{IC} + w_{IV}\mathcal{L}_{IV}$$

 $\mathcal{L}^{data}$

measurements
optional

 $\mathcal{L}_p^{d_e}$

PDE residual
always present

 $\mathcal{L}^{DBc}$

Dirichlet
if not hard-enforced

 $\mathcal{L}^{NBc}$

Neumann
if flux prescribed

 $\mathcal{L}^{Ic}$

initial state
time-dependent

 $\mathcal{L}^{Iv}$

initial velocity
2nd order in t

- every term is a mean-squared residual over its own sample set
- so all six are non-negative, and vanish together on the exact solution
- **the PDE term is always present** — which others appear is set by the problem


The PDE Residual Term

COLLOCATION FORM

$$\mathcal{L}_{PDE} = \frac{1}{n_{PDE}} \sum_{i=1}^{n_{PDE}} [\mathcal{F}_{\hat{w}}(\mathbf{x}_i, t_i)]^2$$

- the residual is what remains when the network is substituted into the equation
- **squared point by point, then averaged** — so the term is non-negative
- the collocation points live strictly inside the domain
- they carry no boundary information whatsoever

RESIDUAL MAGNITUDE OVER THE DOMAIN Ω



Reading this figure. The axes are the spatial domain Ω — a map of where the network sits, not a loss curve. Each dot is one collocation point, and its size and colour encode $|\mathcal{F}_{\hat{w}}|$: how much the network violates the PDE there.

Dirichlet Conditions: Soft Enforcement

BOUNDARY RESIDUAL

$$\mathcal{L}_{DBC} = \frac{1}{n_{DBC}} \sum_{i=1}^{n_{DBC}} [u_{\hat{w}}(\mathbf{x}_i, t_i) - \bar{u}(\mathbf{x}_i, t_i)]^2$$


- the prescribed value is compared with the output there
- soft enforcement means only approximately satisfied
- it competes with the PDE residual for the same parameters
- that competition is why a PINN ignores its boundary

PHYSICAL EXAMPLE — A LIQUID-COOLED POWER MODULE

The baseplate sits on a coolant channel held at 40 °C. That is a Dirichlet condition — the temperature on that face is prescribed, not computed.

BOUNDARY SAMPLING

interior points (PDE only)



$u = \bar{u}$ enforced at these points

Hard Enforcement: the Trial Solution

TRIAL SOLUTION ON A RECTANGLE $[A,B] \times [C,D]$

$$u_{\text{trial}}(x, y) = x(1-x)y(1-y) \mathcal{N}_{\widehat{w}}(x, y)$$

Trial solutions used in the book

- the polynomial factor vanishes on all four edges
- so the trial solution satisfies the condition by construction
- **the boundary loss term disappears entirely** — nothing to balance
- training is markedly easier, and the search space is smaller
- **the cost is generality** — the factor must be built per geometry

The factor is chosen per geometry; Chapter 3 trains a small network to represent it on irregular domains.

$$f_{DBC} = x(1-x)y(1-y)$$

zero on all four edges

$$f_{DBC} = x$$

zero on the left edge only

$$f_{DBC} = \sin(\pi x)\sin(\pi y)$$

trigonometric alternative

Hard vs. Soft Enforcement, Measured

BENCHMARK PROBLEM

$$\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} + 2\pi^2 \sin(\pi x) \sin(\pi y) = 0$$

on the unit square, $u = 0$ on all four edges
 $u(x, y) = \sin(\pi x) \sin(\pi y)$

What this equation is. Poisson’s equation on the unit square with a manufactured source term, so the exact solution is known in closed form and every error below is measured against truth rather than estimated.

tanh network, 40 × 2 | P = 85 | 320 collocation + 240 boundary points | Adam then L-BFGS, 1000 iterations*

METRIC	SOFT BC	HARD BC
Relative L ₂ error	2.12e-02	1.29e-04
Max absolute error	5.14e-02	2.73e-04
Max error on boundary	5.14e-02	0.00e+00
PDF residual MSE	8.47e-03	4.55e-04

Hard enforcement is ~165× more accurate in relative L₂, at essentially the same cost.

WHAT THE NUMBERS CONFIRM

The soft-BC model's largest error sits on the boundary it was penalised to satisfy: an approximation, not a guarantee.

The hard-BC trial solution makes it identically zero — measured, not asserted.


Neumann Conditions and the Normal Derivative

NEUMANN RESIDUAL

$$\mathcal{L}_{NBC} = \frac{1}{n_{NBC}} \sum_{i=1}^{n_{NBC}} [\mathcal{B}_{\hat{u}}(\mathbf{x}_i, t_i)]^2$$

- $\mathcal{B}$ constrains a derivative on the boundary, not a value
- flux, traction, insulation
- the gradient comes from autograd, exactly as for the interior

the outward normal, at every boundary point


 $\nabla \hat{u}$

autograd gives the gradient

 $\mathbf{n}$

the geometry gives the normal

 $\nabla \hat{u} \cdot \mathbf{n}$


their dot product is the flux

$\mathcal{B}$ constrains that number, not $\hat{u}$ itself

NORMAL DERIVATIVE

$$\frac{\partial u}{\partial n} = \nabla u \cdot \mathbf{n}$$

OUTWARD NORMAL ON $\partial\Omega$



Initial Conditions and Initial Velocity

INITIAL CONDITION


$$\mathcal{L}_{IC} = \frac{1}{n_{IC}} \sum_{i=1}^{n_{IC}} [\mathcal{I}_{\hat{w}}(\mathbf{x}_i, t_i)]^2$$

INITIAL VELOCITY

$$\mathcal{L}_{IV} = \frac{1}{n_{IV}} \sum_{i=1}^{n_{IV}} [\dot{\mathcal{I}}_{\hat{w}}(\mathbf{x}_i, t_i) - \dot{\mathcal{U}}(\mathbf{x}_i, t_i)]^2$$

- both terms are sampled over the domain at $t = 0$
- **first order in time** — heat — needs one initial condition
- **second order** — wave — needs the initial velocity too
- without the initial velocity the problem is under-determined
- **the dotted quantities come from autograd** — no separate discretisation

SPACE-TIME DOMAIN




Weighting the Loss Terms

- the weights are hyperparameters; the terms differ in units
- so the raw magnitudes are not comparable
- when one term dominates, the optimiser ignores the others
- fixed weights are cheap, problem-specific and fragile
- **Liu's position** — restructure so the awkward term vanishes

GRADIENT-NORM BALANCING

$$w_k \leftarrow \frac{\|\nabla_{\tilde{w}} \mathcal{L}_{PDE}\|}{\|\nabla_{\tilde{w}} \mathcal{L}_k\|}$$

GRADIENT MAGNITUDE PER TERM



Before balancing, the boundary and initial terms contribute almost nothing to the parameter update. After rescaling, every constraint influences training at a comparable rate.

Three Ways to Evaluate the Residual

STRONG FORM

1

Collocation

Residual enforced pointwise, derivatives from autograd.
Mesh-free; needs the highest differentiability of σ .
The PINN default.

STRONG FORM, STENCIL

2

Finite difference

Do not differentiate the network at all. Evaluate it at neighbouring points a distance h apart and take differences — the classical FDM formula on network outputs.

WEAK FORM

3

Energy method


Minimise a potential-energy functional instead of a residual. Only first derivatives appear. Requires quadrature weights and essential BCs to be hard-enforced.

- all three minimise the same physics
- they differ in how much differentiation is pushed onto the network
- the weak form trades derivative order for integration
- lowering what the activation must supply, at the price of quadrature

Collocation: Sampling the Strong Form


- the residual is enforced at a finite set of points
- **between them nothing is imposed** — so density limits accuracy

UNIFORM GRID




predictable, but density is uniform where it need not be

LATIN HYPERCUBE



one point per row and column:
better coverage, fewer points

RESIDUAL-ADAPTIVE



points follow the residual, and
concentrate at sharp features


Finite-Difference Residuals (Central Scheme)

SECOND-ORDER CENTRAL DIFFERENCES

$$u_{,xx} \approx \frac{u(x-h, y) - 2u(x, y) + u(x+h, y)}{h^2}, \quad u_{,yy} \approx \frac{u(x, y-h) - 2u(x, y) + u(x, y+h)}{h^2}$$

- a trained network can be queried anywhere, so stencils apply
- the Laplacian is assembled from four neighbouring evaluations
- the five network evaluations are each differentiable in the parameters
- **so training still works** — the graph just stays shallow
- **the trade is explicit** — truncation error returns, and h must be chosen

FIVE-POINT STENCIL



four extra network evaluations per collocation point

The Energy Method (Weak Form)

TOTAL POTENTIAL ENERGY FOR POISSON'S EQUATION

$$U_T(u) = \frac{1}{2} \int_{\Omega} \nabla u \cdot \nabla u \, dx - \int_{\Omega} u f \, dx - \int_{\partial\Omega_{NBC}} u t \, ds$$

STATIONARITY

$$\delta[U_T(u)] = \delta[U - W_f - W_t] = 0$$

Strong form

$$-\Delta u = f$$

Test function

multiply by v , integrate

Integrate by parts

order reduced to first

Energy functional

$$U^T = U - W_f - W_t$$

Minimise

$$\delta U^T = 0$$

- integration by parts moves one derivative onto the test function
- **so only first derivatives of u remain** — a weaker demand on the activation
- **the Dirichlet condition becomes essential** — it must hold for the principle to apply
- Neumann conditions appear as a boundary work term and need no loss

T3 Triangulation and Integration Weights

ELEMENT AREA

$$A = \frac{1}{2} |\mathbf{u} \times \mathbf{v}| = \frac{1}{2} |u_x v_y - u_y v_x|$$

NODAL WEIGHT


$$w_i = \sum_{T \ni i} \frac{A_T}{3}$$

QUADRATURE


$$\int_{\Omega} g \, d\Omega \approx \sum_{i=1}^n w_i g(\mathbf{x}_i), \quad \sum_{i=1}^n w_i = A_{\Omega}$$

- the domain is triangulated into three-node cells
- the mesh does not store unknowns** — the network is still the solution
- each triangle gives a third of its area to each vertex

three-node cells over the domain



one node, and the cells around it




each gives it a third of its area

the weights sum to the domain area, so the quadrature is consistent on any mesh

94 triangles, 50 nodes

T3 MESH WITH NODAL WEIGHTS



marker size $\propto$ nodal weight w_i

The Training Loop

1 Sample

Draw interior, boundary and initial points; attach weights if the residual is integral.

2 Forward

Evaluate $\mathbb{N}\hat{w}$ at every sample point in one batched pass.

3 Differentiate

Autograd produces all spatial and temporal derivatives required by $\mathcal{F}$, $\mathcal{B}$, $\mathcal{I}$ and $\mathcal{V}$.

4 Assemble

Form each squared residual and combine into the composite loss $\mathcal{L}$.

5 Update

Backpropagate $\partial\mathcal{L}/\partial\hat{w}$ and step the optimiser: Adam first, then L-BFGS.

6 Adapt

Optionally resample where the residual is largest, then repeat until the criterion is met.

Only step 5 touches the parameters. Everything else is the construction of a scalar functional — which is why the same loop trains every PDE in L7.2 unchanged.


Two-Stage Optimisation: Adam then L-BFGS

- ▶ Adam is first-order and robust to poor initialisation, but it stalls at moderate residual.
- ▶ L-BFGS is quasi-Newton: fast near a minimum, but it diverges if started far from one.
- ▶ Running Adam first and L-BFGS second exploits both.
- ▶ Liu's schedule: Adam at $\text{lr} = 1\text{e-}3$ for ~ 201 epochs, then L-BFGS at $\text{lr} = 1$ with strong-Wolfe for ~ 11 . An L-BFGS epoch costs $55\text{--}750\times$ an Adam epoch.
- ▶ Adam alone is not merely slower: past $\sim 2\,500$ epochs it oscillates, and its loss can get worse.

MEASURED — SAME PROBLEM, NETWORK AND SEED

Two-stage	6.5e-06	18 s
Adam alone	4.4e-04	126 s

LOSS HISTORY (LOG SCALE)



switch at the plateau, not at a fixed loss value

Optimisers in Context: Why L-BFGS Works Here

- ▶ The two-stage recipe looks unusual against mainstream deep learning, where second-order methods are almost never used. The reason it works for PINNs is structural, not empirical.

SGD

the classical baseline

Momentum variants can find minima that generalise better (Wilson et al., 2017).
Slow on ill-conditioned losses.

Adam

adaptive moments

With tuned hyperparameters it matches SGD and converges faster (Choi et al., 2019).
Liu's stage 1, at $\text{lr} = 1\text{e-}3$.

L-BFGS

quasi-second-order

Rare in mainstream training — it needs a deterministic, full-batch loss.
A PINN supplies exactly that.

Why the PINN recipe differs: mainstream training is stochastic and minibatched, so curvature estimates are unreliable. A PINN's collocation set is fixed and its loss deterministic — which makes L-BFGS usable, and it is what delivers the final orders of magnitude.

Adaptive Sampling and Residual Normalisation

RESIDUAL-ADAPTIVE REFINEMENT

$$\mathbf{x}_{new} = \arg \max_{\mathbf{x} \in S} |\mathcal{F}_{\hat{w}}(\mathbf{x})|$$

- Adaptive refinement adds points where the residual is largest
- reaching an accuracy with far fewer points than uniform
- normalisation rescales each term, so none dominates on units
- both are stabilisers, not fixes

WHEN IT BECOMES NECESSARY, NOT MERELY HELPFUL

Burgers' equation at $\nu = 0.01/\pi$ develops a near-discontinuity by $t \approx 0.4$.
Uniform sampling spends its points where the solution is flat.


RESIDUAL NORMALISATION

$$\tilde{\mathcal{L}}_k = \frac{\mathcal{L}_k}{\mathcal{L}_k^{(0)} + \epsilon}$$

BEFORE / AFTER REFINEMENT

initial

refined



dashed line: region of large $|\mathcal{F}_{\hat{w}}|$


Spectral Bias: Low Frequencies Arrive First

TARGET USED IN THE DEMONSTRATION

$$u(x) = \sin(2\pi x) + \frac{1}{2}\sin(10\pi x)$$

- ▶ Gradient descent recovers the smooth part of a target long before its fine structure. Fitting the target above, the low-frequency mode reaches 90% of its amplitude by step 425; the high-frequency mode is at 40% after 6 000 steps.
- ▶ This is the F-Principle: descent converges fastest along the low frequencies.
- ▶ The consequence is subtle — the loss can be small and still falling with the fine structure absent.
- ▶ Mitigations: Fourier feature embeddings, SIREN activations, multi-scale or hierarchical architectures.

AMPLITUDE RECOVERED, AS A FRACTION OF TARGET



Training step →. The gap between the two curves never closes on any budget a student will run.

Failure Modes Worth Recognising

GRADIENT IMBALANCE

PDE-residual gradients overwhelm the BC and IC terms, so the network solves the equation while ignoring the conditions that make the answer unique.

Symptom: total loss falls, boundary error does not. Check per-term curves.

CAUSALITY VIOLATION

Nothing in the plain loss enforces time ordering, so training can fit late times before early ones have converged.

Fix: weight the time terms so the solution converges outward from $t = 0$.

STIFF AND MULTI-SCALE PDES

Widely separated scales produce an ill-conditioned optimisation landscape that traps both Adam and L-BFGS.

Symptom: oscillation near the minimum rather than convergence.

SENSITIVITY TO SETUP

Results shift with the collocation distribution and with weight initialisation — two runs of the same model need not agree.

Fix: fix the seed and report it, as Liu does throughout (seed 88).

Every item here is diagnosable from plots you already produce — per-term loss curves and a residual heatmap. Check the terms separately before changing the architecture.

Key Takeaways

- 1 The ansatz** A PINN replaces the unknown field with a differentiable network. The Universal Approximation Theorem guarantees such a representation exists; it guarantees nothing about finding it.
- 2 The operator** Autograd supplies exact derivatives of that network, so the PDE, boundary and initial operators can all be evaluated pointwise without any discretisation.
- 3 The functional** Physics enters as a composite loss of squared residuals. $\mathcal{L}_p^d$ is always present; the remaining terms depend on the PDE class and on how much is hard-enforced.
- 4 The trade** Hard enforcement removes a loss term and the balancing problem with it — at the cost of a geometry-specific construction. This is the recurring design decision in L7.2.
- 5 The discipline** Capacity and sampling are coupled through $N_{\text{sample}} \geq P^*$, and optimisation is two-stage by default: Adam to explore, L-BFGS to converge.

Next — L7.2: the fundamental PDEs (Poisson, wave, Helmholtz, Burgers, transient heat) behind L8 heat, L9 fluids and batteries, L10 robotics and L11 power grids.

PINNs at a Glance: Strengths, Limits, Best Fit

STRENGTH

Mesh-free and differentiable end to end. Fuses sparse data with physics. Returns a continuous solution evaluable at any point, with no interpolation.

WEAKNESS

Spectral bias and gradient pathologies. Training cost far above a comparable classical solve for a single forward problem.

BEST FIT

Sparse-data regimes, inverse and parameter-identification problems, smooth to moderately oscillatory PDEs, repeated parametric studies.

CAUTION

Stiff, high-frequency and multi-scale problems need special treatment — or a classical solver. Judge the method against the problem, not the fashion.

Physics enters through the loss, not the architecture — which is why the same network, unchanged, solves heat conduction in L8 and a battery thermal model in a project.

Questions

TEN TO TEST WHETHER TODAY LANDED

- **what does the PDE residual measure** — and where is it evaluated?
- why must the activation be smooth, stated from the order of the equation?
- what does universal approximation not promise?
- **soft or hard enforcement** — what does each cost you?
- why do the collocation points carry no boundary information?
- what is the weight between two loss terms, physically?

THE ONE WITH NO SETTLED ANSWER

There is no validation set here. How will you know when to stop?

BEFORE L7.2

Ex_07.1 notebooks 01 and 02. Bring both error figures.


Ex_07 — PINN for the 2-D Poisson Problem

STUDENT TASKS

- ① Implement the PINN; generate interior and boundary collocation points.
- ② Train a soft-BC model; record the final PDE and BC losses separately.
- ③ Train a hard-BC model using the trial solution.
- ④ Compare both against the analytical solution on a regular grid.
- ⑤ Report relative L_2 error and maximum absolute error.

QUESTIONS — AND WHERE THIS DECK ANSWERS THEM

- | | |
|--|---------------|
| Why is a smooth activation preferred? | slide 11 |
| What is the role of automatic differentiation? | slide 14 |
| What does hard enforcement gain over a weighted BC loss? | slides 22, 23 |
| PDE loss small but boundary inaccurate — what to check? | slides 26, 36 |
| How do collocation count and spread affect the solution? | slides 9, 28 |



Measured reference results for both variants are on slide 23 — attempt the exercise before reading them.

Source Map and Merge Log

- ▶ This deck merges the mathematical spine of the L7.1 rewrite with the concrete measurements, course framing and exercise material from the L7_PINN deck.

L7.1 rewrite (28 slides)	spine — five-stage flow, typeset equations, diagrams
L7_PINN (43 slides)	merged in — measured benchmarks, course arc, exercise, limits
Liu (2025), Ch. 2–6	primary source for all formulae and logged numbers
Raissi et al. (2019)	operator form, inverse problem, Burgers benchmark

ADDED IN THIS MERGE

Course arc L8–L12 · modelling spectrum · measured hard-vs-soft benchmark · Liu's real network configurations · Adam/L-BFGS logged numbers · optimiser context · spectral bias · failure modes · strengths and limits · Ex_07 exercise

SOURCE CORRECTIONS CARRIED FORWARD

P^* is the pseudo-dimension $(p+1) + (N_n+1)N_I$, not the trainable-parameter count — both source decks are corrected here.

Chapter pointers: the samples–neurons theorem is Liu Ch. 2.5.2 Eq. (2.19), not Ch. 2.3; `train_adam_lbfgs` is Ch. 2.7.5–2.7.6, not Ch. 2.4.