

The image features the text "DL for ENG" in a bold, 3D, gold-colored font. The letters are floating on a deep blue ocean with visible ripples and small waves. In the background, there are rolling blue hills or mountains under a sky filled with soft, grey clouds. The overall scene has a serene, digital aesthetic.

DL for ENG

LECTURE 8.1

Stationary Heat


Remus Teodorescu

*With support from Research Assistant Noman Khan -
nomank@energy.aau.dk*

Aalborg University · ret@et.aau.dk

Deep Learning for Engineering · 2026

CONDUCTION · GEOMETRY · THE PLATE



Steady conduction is Poisson's equation with units: Fourier's law, a conductivity, and walls held at a temperature, a flux or a convection. On a plate with a cooling hole the hole is built into the network by its level set, and the answer is checked by an energy balance.

Learning Objectives

SECTIONS, IN ORDER

- 1 **Conduction** — Fourier's law, variable k , the thermal walls, scaling
- 2 **Geometry** — sampling a domain with a hole, its multiplier, its normal
- 3 **Plate with a hole** — verify, solve, weight, and balance the heat

READ, AND WHAT EACH ANSWERS

- | | |
|-------------|--|
| Liu ch. 2.4 | what does a real thermal problem need? |
| Liu ch. 3 | how is a domain with a hole handled? |
| Liu ch. 4.6 | how is the answer checked? |

Recap

1 – 2 · STEADY

1 What decides the type of a second-order equation, and which conditions does each type need?

How it treats time. An elliptic equation has none: a condition on every wall. A parabolic one has one time derivative: the walls and one initial field. A hyperbolic one has two: the walls and two initial conditions. With fewer, the problem has many answers. *Liu ch. 2; Ex_07.2 notebooks 1 – 3*

2 What do Poisson's and Helmholtz's equations describe, and what must a Helmholtz network be protected from?

Poisson is a steady balance between a source and diffusion. Helmholtz adds a term in the field itself, so it has modes, and $u \equiv 0$ always satisfies it: the network needs a unit-norm term and a floor on λ . In the slot, λ reaching the first eigenvalue is thermal runaway. *Liu ch. 2.4, 6.2; Ex_07.2 notebook 2*

3 – 4 · IN TIME

3 What does time add: how do the heat and the wave equation treat a disturbance, and an error made early?

Heat spreads and forgets: every feature decays, the sharpest fastest, and an early error is damped. A wave travels at a finite speed and keeps its shape, so an early error is carried along. Heat needs one initial condition, the wave two. *Liu ch. 4.1, 5.1; Ex_07.2 notebooks 1, 3*

4 How does a PINN meet the conditions of each type, and which type is hardest for it?

A condition with a formula is built into the trial function: the walls, the initial field. One without, like the panel's initial velocity, is a term in the loss. The wave is hardest: its error is not damped, and its residual must be scaled before the network moves. *Liu ch. 4.2, 5.2; Ex_07.2 notebook 3*

Steady Conduction

Steady conduction — heat flows down the temperature gradient, $\mathbf{q} = -k\nabla T$ (Fourier's law), and in steady state what leaves a small volume equals what is generated in it, $\nabla \cdot \mathbf{q} = Q$. Together they give Poisson's equation for the temperature, with Q/k as its source.

FOURIER'S LAW IN AN ENERGY BALANCE

$$\mathbf{q} = -k \nabla T, \quad \nabla \cdot \mathbf{q} = Q$$


$$T_{xx} + T_{yy} + \frac{Q}{k} = 0$$

$\mathbf{q}$ the heat flux, k the conductivity in W/(m K)
 Q the heat generated in W/m³

CONDUCTION IN PYTORCH

```
T = net(xy)                # the temperature, in kelvin
q = -k * grad(T, xy)        # Fourier's law; grad = autograd
F = div(q, xy) - Q          # the energy balance, zero inside
```

A WALL THAT GENERATES HEAT



Both faces held at T_0 . The profile is a parabola, and at the amber peak the gradient, and so the flux, is zero: heat leaves through both faces.

Variable Conductivity

Variable conductivity — when k changes with position it stays inside the divergence: $\nabla \cdot (k \nabla T) + Q = 0$. Across an interface between two materials the gradient jumps and the normal flux does not, so the divergence must not be expanded; autograd differentiates $k \nabla T$ as written.

K INSIDE THE DIVERGENCE

$$\nabla \cdot (k(\mathbf{x}) \nabla T) + Q(\mathbf{x}) = 0$$


$k(\mathbf{x})$ the conductivity at each point, $k \nabla T$ differentiated as one quantity

- **at an interface** — the slope jumps, the flux does not
- **in code** — differentiate the product, not its parts

VARIABLE k IN PYTORCH

```
k = k_field(xy)           # the conductivity at each point
q = -k * grad(T, xy)       # Fourier's flux, one quantity
F = div(q, xy) - Q         # differentiated as written
```

TWO MATERIALS, ONE INTERFACE



Left face at 1, right face at 0, no generation. The slope falls fourfold at the interface; the flux $k dT/dx$ is the same on both sides.

Thermal Boundary Conditions

Thermal boundary conditions — three ways to fix a wall: a temperature (Dirichlet), a heat flux (Neumann), or convection to a fluid (Robin), which ties the flux to the wall temperature. Only a temperature can be built into the network; the other two stay loss terms.

A TEMPERATURE, A FLUX, A CONVECTION

Dirichlet

$$T = \bar{T} \text{ on } \partial\Omega_D$$

Neumann

$$-k \frac{\partial T}{\partial n} = \bar{q} \text{ on } \partial\Omega_N$$

Robin

$$-k \frac{\partial T}{\partial n} = h (T - T_\infty)$$

$\bar{T}$ the wall temperature, $\bar{q}$ the wall heat flux, n the outward normal
 h the heat transfer coefficient, T_∞ the fluid temperature

THE WALLS IN PYTORCH

```
B_D = T(x_D) - T_bar           # Dirichlet: the temperature
B_N = -k * dTdn(x_N) - q_bar    # Neumann: Fourier's flux, given
B_R = -k * dTdn(x_R) - h * (T(x_R) - T_inf) # Robin: Fourier's flux, tied to T
```

ROBIN · CONVECTION AT THE WALL

- **balance** — heat conducted to the wall equals heat the fluid carries away
- h — heat transfer coefficient, $\text{W}/(\text{m}^2 \text{K})$, or film coefficient: the fluid layer at the wall in one number
- **typical h** — still air 2 – 25, forced air 25 – 250, forced water 100 – 20 000

BIOT NUMBER

$$\text{Bi} = \frac{hL}{k} = \frac{L/k}{1/h}$$

L the size of the body, L/k resists conduction inside
 $1/h$ resists convection at the wall

$\text{Bi} \ll 1$: nearly insulated (Neumann). $\text{Bi} \gg 1$: the wall is at T_∞ (Dirichlet). 1 mm of silicon in still air: 1.67×10^{-4} .

Well-Posedness

Compatibility condition — a problem with a flux condition on every wall has a solution only if the heat out through the walls equals the heat generated, and then only up to a constant: any shift of T fits the same equation and walls. A PINN picks one without warning.

HEAT OUT EQUALS HEAT GENERATED


$$\oint_{\partial\Omega} \bar{q} ds = \int_{\Omega} Q d\Omega$$

$\bar{q}$ the flux out through the boundary $\partial\Omega$, Q the heat generated inside

- **one value pinned** — fixes the constant
- **one wall at a temperature** — fixes it too

TALK · 2 MIN How would you explain why a problem with only flux conditions has no single solution, in your own words? Then, in the plot on the right, which of the three curves is right, and what would you add to the loss to choose it?

THE SAME FIELD, SHIFTED



FIXING THE LEVEL

```
L_pin = (T(x_ref) - T_ref).pow(2)    # pin one value
L_mean = T(xy).mean().pow(2)         # or fix the mean
```

Non-Dimensional Form

Non-dimensional form — the problem rewritten in scaled variables: lengths divided by the device size L , temperature by a rise $\Delta T = QL^2/k$. Every term is then of order one, as the initialisation and the loss weights assume. Train scaled, and convert back at the end.

SCALE, TRAIN, CONVERT BACK

$$\theta = \frac{T - T_0}{\Delta T}, \quad \xi = \frac{x}{L}, \quad \Delta T = \frac{QL^2}{k}$$

θ the scaled temperature, ξ the scaled coordinate, ΔT the temperature scale

$$\theta_{\xi\xi} + \theta_{\eta\eta} + 1 = 0$$

$\eta = y/L$; the same equation, every coefficient equal to one

SCALING IN PYTORCH

```
dT = Q * L**2 / k           # the temperature scale, kelvin
theta = net(xy_m / L)       # train on scaled coordinates
T = T0 + dT * theta         # convert back, in kelvin
```

A SILICON DEVICE, BEFORE AND AFTER

	physical	scaled
length L	1 mm	1
conductivity k	150 W/(m K)	1
source Q	10^8 W/m ³	1
rise ΔT	$QL^2/k = 0.67$ K	1
coefficient h	25 W/(m ² K)	$Bi = 1.67 \times 10^{-4}$

In metres and kelvin the source term is ten to the eighth and the length ten to the minus third. Scaled, every term is of order one.

Sampling a Real Domain

Rejection sampling — points drawn uniformly over a box around the domain, keeping those the level set places inside, $\phi > 0$. On a curved wall the points are spaced by arc length rather than by angle, so the ends of the hole, where it bends most, get their share.

ONE FORMULA: INSIDE, OUTSIDE, ON THE WALL

$$\phi(x, y) = \left(\frac{x - x_c}{a}\right)^2 + \left(\frac{y - y_c}{b}\right)^2 - 1$$



(x_c, y_c) the centre, a, b the half axes; $\phi < 0$ in the hole, $\phi = 0$ on its wall

- **inside** — keep a point where $\phi > 0$
- **on the wall** — equal steps of arc length


SAMPLING IN PYTORCH

```
xy = torch.rand(4000, 2)           # uniform over the square
xy = xy[phi(xy) > 0][:2000]        # keep the points in the plate
xy_h = ellipse_by_arc_length(160)  # the hole wall, evenly spaced
```

THE PLATE, AND ITS WALL POINTS



equal angle: ends bare



equal arc length

Boundary Multiplier

Boundary multiplier — a function that is zero on a boundary and positive in the domain. Multiplied into the network's output, it makes a value hold there exactly. For an ellipse the level set ϕ is one; for a shape with no formula, a small second network is trained to be one, then frozen.

THE HOLE WALL, BUILT IN


$$T(\mathbf{x}) = \phi(\mathbf{x}) \mathcal{N}_{\hat{w}}(\mathbf{x})$$

ϕ the level set, zero on the hole; $\mathcal{N}_{\hat{w}}$ the network; $T = 0$ there for any weights

- **a value** — held exactly by a factor
- **a derivative** — cannot be, stays a loss term

TALK · 2 MIN How would you explain what a boundary multiplier does, in your own words? Then, in the figure on the right, where is the multiplier small, and what does that do to the network's gradient near the hole?

ϕ OVER THE PLATE



dark: the hole

THE HOLE IN PYTORCH

```
T = hole_multiplier(xy) * net(xy)  # zero on the hole
loss = F.pow(2).mean() + w_N * L_N  # edges still a term
```

Flux on a Curved Boundary

Normal derivative — the rate of change of T across a wall, $\partial T / \partial n = \nabla T \cdot \mathbf{n}$, where $\mathbf{n}$ is the unit normal pointing out of the material. On a curved wall $\mathbf{n}$ changes with position: it comes from the gradient of the level set, and its sign decides whether heat leaves or enters.

THE NORMAL FROM THE LEVEL SET

$$\frac{\partial T}{\partial n} = \nabla T \cdot \mathbf{n}, \quad \mathbf{n} = -\frac{\nabla \phi}{|\nabla \phi|}$$


$\mathbf{n}$ the unit normal out of the plate; minus, because out of the plate is into the hole

- **on an edge** — the normal is a constant
- **on the hole** — it turns with the wall

THE NORMAL IN PYTORCH

```
g = grad(phi(xy_h), xy_h)           # the level set's gradient
n = -g / g.norm(dim=1, keepdim=True) # unit normal, out of the plate
q_n = -k * (grad(T, xy_h) * n).sum(1) # Fourier's flux through the wall
```

NORMALS ON THE HOLE WALL



Each white stick is $\mathbf{n}$ at an amber point: it turns with the wall and always points out of the plate. A sign error inverts the flux, not the magnitude.

Manufactured Solution

Manufactured solution — a solution chosen in advance, with the source derived from it by differentiation, so the answer is known exactly. Solving that problem checks the residual, the trial solution and the training loop before they meet a problem with no answer to compare against.

CHOOSE THE ANSWER, DERIVE THE SOURCE

$$T = \sin(\pi x)\sin(\pi y) \Rightarrow \frac{Q}{k} = 2\pi^2 \sin(\pi x)\sin(\pi y)$$

T the chosen answer, zero on the four edges; Q/k the source derived from it


$$\varepsilon = \frac{\|T_{\hat{w}} - T_{\text{exact}}\|_2}{\|T_{\text{exact}}\|_2}$$

$T_{\hat{w}}$ the network's answer, T_{exact} the chosen one

THE CHECK IN PYTORCH

```
Q_k = 2 * pi**2 * sin(pi*x) * sin(pi*y)    # the source, from the answer
T = x*(1-x) * y*(1-y) * net(xy)            # the walls held exactly
err = (T - T_exact).norm() / T_exact.norm() # compare with the answer
```

FOUR STEPS BEFORE THE PLATE



what it checks

the residual, the trial, the loop

what it cannot

the geometry of the plate

An error found here costs minutes; the same error found on the plate costs an afternoon.

Heated Plate with a Hole

Heated plate with a hole — the case study of this block: a square plate with uniform heat generation, cooled through an elliptical hole held at $T = 0$, with insulated outer edges, $\partial T / \partial n = 0$. The hole wall is built into the network by its level set; the edges stay a loss term.

ONE BUILT-IN WALL, ONE LOSS TERM

$$\mathcal{L} = \mathcal{L}_{PDE} + w_N \mathcal{L}_{NBC}$$


$\mathcal{L}_{PDE}$ the residual inside, $\mathcal{L}_{NBC}$ the insulated edges, w_N their weight

- **the hole** — $T = 0$, by the multiplier
- **the edges** — no flux, a weighted term

THE PLATE IN PYTORCH

```
T = hole_multiplier(xy) * net(xy)          # the hole held at zero
L_N = dTdn(T, xy_edges).pow(2).mean()      # the edges: no flux
loss = (lap(T) + Q_k).pow(2).mean() + w_N * L_N # residual and edges
```

THE PLATE OF EX_08.1



amber wall

$T = 0$, built in

dashed edges

$\partial T / \partial n = 0$, a loss term

Flux Weight

Flux weight — the weight w_N on the insulated-edge term. Too small, and the network ignores the edges and heat leaks through them; too large, and the residual inside pays for a perfect wall. It is chosen by a sweep that watches both terms, not the total loss.

ONE RUN PER WEIGHT, BOTH TERMS WATCHED

$$\mathcal{L} = \mathcal{L}_{PDE} + w_N \mathcal{L}_{NBC}$$

$\mathcal{L}_{PDE}$ the residual inside, $\mathcal{L}_{NBC}$ the insulated edges, w_N their weight

- **the total loss** — hides which term is paying
- **the choice** — a trade, argued from both

THE SWEEP IN PYTORCH

```
for w_N in [0.1, 1, 10, 100]:
    net = train(w_N)
    log(w_N, L_pde(net), L_N(net))
```

one run per weight
the same seed, the same points
both terms, not their sum

WHAT EACH END OF THE SWEEP DOES

w_N small

the edges leak heat

w_N large

the residual inside pays

between

argued from both terms

The weight is a rate of exchange between a flux and a residual, which have different units. Scaling changes the rate before any sweep does.

Energy Balance Check

Energy balance check — a check on the trained network, not a new law. Fourier's law gives the heat flux at each point of the hole's wall; added up around the hole, that is the heat leaving the plate. In steady state it equals the heat generated: the energy balance, for the whole plate.

FOURIER'S FLUX, ADDED UP AROUND THE HOLE

$$P_h \overline{q_n} = Q A_\Omega$$

P_h the hole's perimeter, $\overline{q_n}$ the mean of Fourier's flux into it, A_Ω the plate's area

- **at a point** — Fourier's law gives the flux into the hole
- **around the hole** — its mean times the perimeter is the heat out
- **in the plate** — Q times the area is the heat generated

THE BALANCE IN PYTORCH

```
q_n = -k * (grad(T, xy_h) * n).sum(1)    # 1: Fourier's flux at each wall point
out = q_n.mean() * perimeter              # 2: heat out; points evenly spaced
gen = Q * (1 - pi * a * b)                # 3: heat generated; compare with out
```

THE CHECK IN NOTEBOOK 03

	value
heat generated	0.94
heat out	0.92
heat out / generated	98.57 %
out short by	1.43 %

The plate with a hole has no exact solution to compare against, so this balance is the check: heat generated is known, heat out is the network's. Notebook 03 calls it the flux balance. Scaled units; the ratio is taken before rounding. Whether 1.43 % is small enough is your decision.

Advanced Topics

RELATED TOPICS LEFT OUT, AND WHERE THEY MATTER

topic	what it adds	read	where it makes a difference
energy method	an energy, first derivatives only	Liu ch. 2.8	elasticity, structures
triangulated domain	points and weights from a mesh	Liu ch. 3.2	any drawn geometry
convergence study	error against points and size	Liu ch. 4.9.3	a design sign-off
<u>diffusion surrogate</u>	trained on solved cases, then steered	UDL ch. 14 – 18	many boundary conditions
PINN or FEM	when each is the better tool	Liu ch. 2.10	one-off steady solves

WHY THEY ARE LEFT OUT

The plate of Ex_08.1 needs none of them: its hole has a formula, its answer can be checked by an energy balance, and it is solved once. A finite-element solver would solve it too; the method earns its place on the inverse problems of L8.2.

Exercise 8.1

THE NOTEBOOKS, IN ORDER

- **00** — the plate, its hole, its samples, its normals
- **01** — the method, on a manufactured solution
- **02** — the plate with its hole built in
- **03** — the flux weight, and the energy balance
- **04** — compares them, and the report

THE QUESTION EACH NOTEBOOK ANSWERS

- 01** is the method right?
- 02** how hot, and where?
- 03** does the heat balance?

REPORT QUESTION, TO CONCLUDE

The manufactured solution checked the method and the energy balance checked the plate. On a design with no exact solution, which checks would you run, and what result would make you distrust the network?

Key Takeaways

STEADY CONDUCTION

what it is Fourier's law in an energy balance

function T for a given k, Q and walls

train, tune scale first; flux walls weighted

vs Poisson units, k inside, three walls

used for devices, plates, heat sinks

REAL GEOMETRY

what it is a domain with no polynomial factor

function ϕ : inside, multiplier, normal

train, tune arc length on the wall; w_N

vs rectangle n and the factor vary with x

used for plates with holes, device sections

Next — L8.2: the same plate, in time.

Questions

1 – 2 · STEADY CONDUCTION

1 What is steady conduction, and which conditions can a wall carry?

Fourier's law in an energy balance: what leaves a small volume equals what is generated in it, which gives Poisson's equation, with k inside the divergence when it varies. A wall has a fixed temperature, a fixed flux or convection. With a flux on every wall the solution is fixed only up to a constant. *Liu ch. 2.2, 2.4; Ex_08.1 notebooks 01 – 02*

2 Why is a device problem scaled before training, and how is the result checked?

In SI units the terms differ by orders of magnitude; scaled, every term is of order one, as the initialisation and the loss weights assume. Check the code on a manufactured solution first, then the result by the energy balance: the heat leaving equals the heat generated. *Liu ch. 4.4, 4.6; Ex_08.1 notebooks 01, 03*

3 – 4 · REAL GEOMETRY

3 How is a domain with a hole sampled, and how is a temperature built into its wall?

Points are drawn over a box and kept where the level set ϕ is positive; on the curved wall they are spaced by arc length. A boundary multiplier, zero on the wall and positive inside, multiplies the network's output, so the wall temperature holds exactly. *Liu ch. 3.1 – 3.3; Ex_08.1 notebooks 00 – 02*

4 How is a flux condition imposed on a curved wall, and how is its weight chosen?

As a loss term on the normal derivative, the gradient of T along the outward unit normal, which comes from the level set. Its weight is found by a sweep that watches the flux term and the residual separately: too small and heat leaks through the wall, too large and the interior pays. *Liu ch. 2.2.3, 3; Ex_08.1 notebook 03*

NEXT

L8.2 — Transient Heat

The same plate, in time: an initial field, causality, and a diffusivity recovered from sensors.

BEFORE THE NEXT SESSION

read	Liu ch. 3, 4.6
run	Ex_08.1 notebooks 00 to 03
bring	the energy balance, and your w_N
think about	which of your walls can be built in

The steady plate is solved and checked. Next it heats up.