

DL for ENG

LECTURE 12.1

Power Grid Stability Estimation

Remus Teodorescu

With support from Research Assistant Noman Khan -
nomank@energy.aau.dk

Aalborg University · ret@et.aau.dk

Deep Learning for Engineering · 2026

GRID · STATE · INERTIA



A grid is a graph with the power flow equations at every bus, and a control room sees only a few noisy meters of it. The lecture estimates the full state with the estimator every utility runs and with a network judged by the physics, with no true state to train on - then lets the grid move: the swing equation, and a plant's inertia recovered from a recorded disturbance.

Learning Objectives

SECTIONS, IN ORDER

- 1 **The grid** — a graph, the power flow equations, what a room sees
- 2 **State estimation** — least squares, and a network with no true state
- 3 **The swing** — the equation of motion, and the inertia from a record

UDL carries the graph network, Liu the inverse fit. The example is Ex_12.1's six-bus grid, representative of eastern Denmark and not a model of it: the line impedances are textbook values, because no operator publishes measured ones, and results that depend on them must say so.

READ, AND WHAT EACH ANSWERS

UDL ch. 13	how is a graph represented?
Liu ch. 4	what can readings determine?
SciPy	how is an ODE shot and fitted?

Recap of L11.2 Dynamics and Energy

1 – 2 • A NETWORK FOR THE CHEAPEST LAP

1 How can a network find the speed profile of the cheapest lap while never passing the grip limit?

The network takes the position along the lap, as six sine and cosine harmonics so that the end of the lap joins its start, and returns the speed there. The loss is the cost of the lap itself, the energy drawn from the battery plus the price on the lap time, differentiated by autograd. The grip limit is built into the last layer: a logistic function holds the speed between the stall speed and the local ceiling for any weights, where a penalty could be traded away for a lower cost. *Liu ch. 5; Ex_11.2 section 5*

2 When does a network for the speed profile pay against the classical optimiser, and when does it not?

In the exercise the classical optimiser, with one bounded speed per step of the lap, and the network agree within 0.05 per cent, and for one lap on a known course the classical optimiser is the right tool. The network pays when its answer is reused: it is a function of the position rather than a list of numbers, its grip limit holds by construction, and it can become a policy that reacts to what the camera sees. *Liu ch. 5; Ex_11.2 sections 4 – 5*

3 – 4 • THE ROLLING COEFFICIENT

3 How can the rolling coefficient of the car be estimated with a network from a coast-down?

The drive is cut and the speed is logged as it falls; with no current, only rolling resistance and air drag slow the car. A small network takes the time and returns the speed, and the rolling coefficient is one more trainable number, kept positive as an exponential. The loss is the residual of the freewheeling equation, computed by autograd, plus the misfit to the logged speed. On the synthetic log it returned the coefficient the log was made with, as least squares did. *Liu ch. 4; Ex_11.2 section 6*


4 What must the measurements contain for the rolling coefficient to be identified, and when is it trusted?

In a coast-down no current flows, so only rolling resistance and drag slow the car, and the decay of the speed separates them; in a steady cruise the rolling resistance and the winding resistance of the motor trade off, and a good fit can hide wrong values. The car has no encoder, so its speed comes from lap timing or the camera. Every car number in the notebook is a placeholder; the model is trusted only after it has predicted a lap that was then measured. *Liu ch. 4; Ex_11.2 section 6 and on the car*

A Grid Is a Graph

Transmission grid — buses joined by lines: a graph, with a node's state being its voltage magnitude and angle, and every edge carrying an admittance. The state of the whole grid is one number pair per bus, and everything a control room wants to know is a function of it.

EX_12.1'S SIX BUSES, AS THEY ARE WIRED



THE CASE, AND ITS LIMITS

the grid 6 buses, 6 lines, 100 MVA base

the state 6 magnitudes, 5 angles

stands for eastern Denmark, loosely

impedances textbook values, flagged

so say when a result leans on them

No operator publishes measured impedances; representative is the honest word.

The Power Flow Equations

Power flow equations — two algebraic balances at every bus: the active and the reactive power the lines bring must equal what the bus injects or takes. They are constraints the voltages satisfy at every instant, not a simulation in time - the elliptic habit of L8.1, on a graph.

TWO BALANCES AT EVERY BUS

$$P_i = V_i \sum_k V_k (G_{ik} \cos \theta_{ik} + B_{ik} \sin \theta_{ik}), \quad Q_i = V_i \sum_k V_k (G_{ik} \sin \theta_{ik} - B_{ik} \cos \theta_{ik})$$

V_i, θ_i the voltage magnitude and angle at bus i : the state, G, B the admittance matrix, $\theta_{ik} = \theta_i - \theta_k$

ONE BUS, ONE BALANCE



HOW THE EXERCISE USES THEM

as a solver	Newton's method: the reference
as a check	the residual at every bus
as a loss	the term that replaces truth
in torch	so autograd reaches the state


The same equations serve three ways, and the third is this lecture's turn: a residual can judge an estimate at a bus no meter sees.

What the Control Room Sees

Measurement model — every reading is a function of the state plus noise: a PMU reads a bus's magnitude and angle, a meter its P and Q , a forecast stands in where no meter exists. The thin set has fewer readings than unknowns, and where they sit on the graph decides everything.

ONE LINE PER METER

$z_j = h_j(V, \theta) + \varepsilon_j$
 z_j one reading, h_j what that meter measures as a function of the state, ε_j its noise



TALK • 2 MIN How would you explain, in your own words, why a grid's state can be known better than its meters? Then, from the table on the right, why does the thin set fail where the forecast set does not?

THE THREE SETS OF EX_12.1

set	readings	covers
full	14	every unknown
thin	6	5 of 11 unknowns
forecast	6 + priors	every unknown, loosely

The forecast's accuracy, 0.05 p.u., is an assumption of the exercise. Observability is a property of where the meters sit on the graph, not of how good they are: the thin set leaves half the state unseen however exact its six readings become.

Weighted Least Squares

State estimation — finding the voltage magnitude and angle at every bus from fewer, noisy readings. The estimator every utility runs is weighted least squares: each meter's misfit, weighted by its accuracy, minimised by Newton's method. It is the baseline any network must face.

THE ESTIMATOR IN SERVICE

$$\hat{x} = \arg \min_x \sum_j \frac{(z_j - h_j(x))^2}{\sigma_j^2}$$

$\hat{x}$ the estimated state, σ_j each meter's accuracy : a good meter weighs more, solved by Newton's method

THE THIN SET: 5 OF 11 UNKNOWNNS



determined by the six readings

unseen, whatever the noise

MEASURED, OVER 300 OPERATING POINTS

set	worst V	angle
full, 14	0.0030 p.u.	0.07°
thin, 6	0.0248 p.u.	1.82°
forecast	0.0050 p.u.	0.15°

One estimate costs about a millisecond. The thin row is not a solver failure: it is the graph saying no.



Judged by the Physics

Estimation with no true state — a network maps the readings to the state, and its loss has no true state in it: the misfit to the meters, plus the power flow residual at every bus. The physics constrains the buses no meter sees, and the same residual is the estimate's own health check.

THE LOSS, WITH NOTHING TO COPY

$$\mathcal{L} = \sum_j \frac{(z_j - h_j(\hat{x}))^2}{\sigma_j^2} + w_P \sum_i r_i(\hat{x})^2, \quad \text{no true state in it}$$

r_i the power flow residual at bus i : physics constrains every bus, metered or not, w_P its weight

THE LOSS IN PYTORCH

```
net = MLP(n_in=14, n_hidden=64, n_layers=3)    # readings in, the state out
V, th = net(z)                                # 9995 weights in all
r = power_flow_residual(V, th)                 # the physics, in torch
L = wls_misfit(z, V, th) + W_P * r.pow(2).sum() # no true state anywhere in it
```

AGAINST THE BASELINE, HONESTLY

accuracy 0.0043 p.u.: it does not win

one estimate 0.26 ms; 2 μ s in a batch

training 13 s, once

a line out, untold both fail alike

but the check its residual: 0.28 to 2.35

The gain is the time of one estimate, and a self-check that says when not to trust it.

The Swing Equation

Swing equation — the grid set in motion: a machine's rotor accelerates with the gap between its mechanical and its electrical power, held back by inertia and damping, and the electrical power ties the machines together. One ordinary differential equation per machine.

THE EQUATION OF MOTION

$$\frac{d\delta}{dt} = \omega - \omega_s, \quad \frac{2H}{\omega_s} \frac{d\omega}{dt} = P_m - P_e(\delta) - D \frac{\omega - \omega_s}{\omega_s}$$

δ the rotor angle, ω its speed, H the inertia in seconds, D damping, P_m, P_e mechanical and electrical power

AFTER A FAULT



the rotor swings, and settles

H sets how slowly

D sets how it dies away

WHY INERTIA IS THE NUMBER

H seconds of stored spin

falling converters replace machines

so the grid swings faster

and operators must know today's H

from what the disturbances it records

DK2 swings with the Nordic area, so the Nordic frequency is the right record for a DK2 plant.

The Inertia From a Record

Inertia identification — the inverse problem of the block: H and D trained beside the weights so the swing equation fits a recorded disturbance, as L8.2's diffusivity was. The residual is written with H on the left, so a large H cannot hide, and only the record after the fault is used.

THE FIT, WRITTEN SO IT CANNOT HIDE

$$H, D \text{ trainable, } \mathcal{L} = \left\| \frac{2H}{\omega_s} \dot{\omega} - P_m + P_e + D \Delta\omega \right\|^2 + w_D \mathcal{L}_{\text{record}}$$

$\mathcal{L}_{\text{record}}$ the misfit to the recorded frequency, after the fault is cleared; H on the left, so a large H cannot hide

THE FIT IN PYTORCH

```
H = torch.nn.Parameter(torch.tensor(8.0))      # trained with the weights
res = (2*H/w_s)*grad(w, t) - P_m + P_e + D*dw    # H on the left
L = res.pow(2).mean() + W_D*(w - rec).pow(2).mean() # after the fault
```

TALK • 2 MIN How would you decide, in your own words, which part of a record identifies a parameter? Then, from the table on the right, why does the window before the fault fail from both starts?

RECORDED WITH AN INERTIA OF 4.00 s

fit	start 8 s	start 2 s
shooting, after	7.41 s	1.83 s
physics-informed, after	3.89 s	3.90 s
shooting, before	50 s, the bound	1.94 s
physics-informed, before	8.35 s	2.10 s

Shooting has local minima on this record; the physics-informed fit did not - a finding about this record and these starts, not a law. Before the fault nothing excites the machine, and no method identifies what the data does not carry.

Advanced Topics

RELATED TOPICS LEFT OUT, AND WHERE THEY MATTER

topic	what it adds	read	where it makes a difference
bad-data detection	a wrong meter, found by its residual	Liu ch. 4	every estimator in service
placing one more meter	where a reading buys observability	pandapower docs	the thin set's missing half
the critical clearing time	how long a fault may stay	PowerSystems.jl	protection settings
rate of change of frequency	the first seconds of a real record	ENTSO-E data	low-inertia grids
the graph network's layers	messages passed along the lines	UDL ch. 13	grids too large for one MLP; L12.2


WHY THEY ARE LEFT OUT

Ex_12.1 estimates one six-bus grid from fourteen readings and fits one machine's inertia, scored against Newton's power flow, so none of them is needed to do it. The last row is where the next lecture begins.

Exercise 12.1

ONE NOTEBOOK, SECTION BY SECTION

- 1 – 3 — the grid, the three sets of readings, the power flow
- 4 — Newton's reference; weighted least squares on each set
- 5 — the network judged by the physics; then a line out, untold
- 6 — the inertia from the recorded disturbance, two ways
- 7 – 9 — the report; mini projects MP12.1A and MP12.1B



THE QUESTION EACH PART ANSWERS

- 1 – 3 what is the state of a grid?
- 4 how far do the meters carry?
- 5 what replaces the true state?
- 6 what does a record identify?

REPORT QUESTION, TO CONCLUDE

The network matched least squares from the same readings and answered three hundred estimates in a blink; its residual rose eightfold when a line went out untold. In a control room, what would you use each tool for?

Key Takeaways

STATE ESTIMATION

what it is	the state from few, noisy meters
function	every bus's V and θ
train, tune	no true state: physics in the loss
vs least squares	equal accuracy, a fraction of the time
used for	a room's picture, and its self-check

SWING DYNAMICS

what it is	the grid's equation of motion
function	angles and speeds after a fault
train, tune	H, D beside the weights
vs shooting	no local minima on this record
used for	today's inertia, from records

Next — L12.2: whether the grid survives the next fault, predicted before it happens.

Questions

1 – 2 • A NETWORK FOR THE STATE

1 How can a network estimate the state of the grid from the readings when no true state is available for training?

The network takes the fourteen readings of the forecast set and returns the voltage magnitude and angle at every bus. From the state it proposes, the power flow equations give what every meter should read, and the loss is the weighted misfit between those values and the actual readings. The reference angle is fixed at zero and the magnitudes are held between 0.90 and 1.10 per unit by construction. No true state appears anywhere in the training. *UDL ch. 13; Ex_12.1 section 5*

2 When does a network for the state of the grid pay against weighted least squares, and how is its estimate checked?

From the same readings it reaches the same accuracy as weighted least squares, so it does not gain in accuracy. Its gain is time: once trained, one estimate takes a quarter of a millisecond, and a few microseconds each in a batch. Both rest on the model of the grid, and with a line out and nobody told both are wrong. The misfit of the network's state to the readings then rises about eightfold, a check that needs no true state. *Liu ch. 4; Ex_12.1 sections 4 – 5*

3 – 4 • THE INERTIA FROM A RECORD

3 How can the inertia of a machine be estimated with a network from a recorded disturbance?

A network takes the time after the fault and returns the rotor angle and speed of the machine, and the inertia is one more trainable number, kept positive as an exponential. The loss is the residual of the swing equation, computed by autograd, plus the misfit to the recorded frequency. The residual is written with the inertia multiplying the acceleration: divided by it, a very large inertia makes the residual small whatever the swing does. *Liu ch. 4; Ex_12.1 section 6*

4 Which part of a record identifies the inertia, and how does the network fit compare with the classical fit?

Only the window after the fault, where the machine swings, carries the inertia. Before the fault nothing moves, and every fit returns a number near its own start or at its bound, with a perfect misfit. After the fault the network fit returned 3.89 s from both starts against the true 4.00 s, while the classical fit by integration stayed near where it was started. That is a finding on this record, and a fit is believed only when two starts agree. *Liu ch. 4; Ex_12.1 section 6*

NEXT

L12.2 — Power Grid Stability Prediction

From what the grid is to what it will do: whether it survives the next fault, predicted by a graph network in the time a protection relay has.

BEFORE THE NEXT SESSION

read	UDL ch. 13
run	Ex_12.1, sections 1 to 6
bring	your inertia, from both starts
think about	what the residual map told you

The estimation half of the block is done: the state, and the machine behind it.